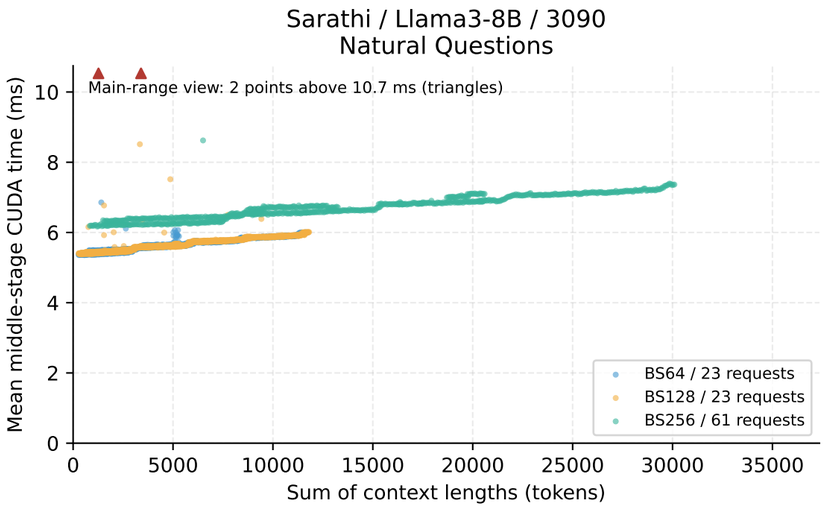
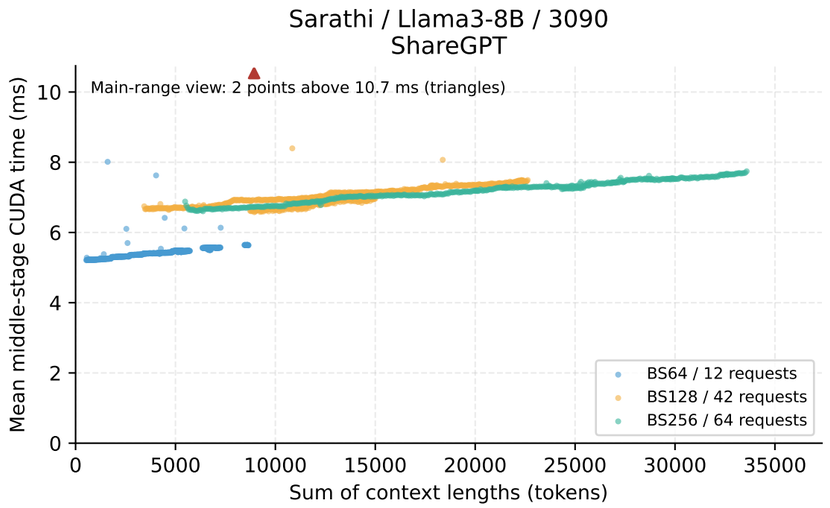
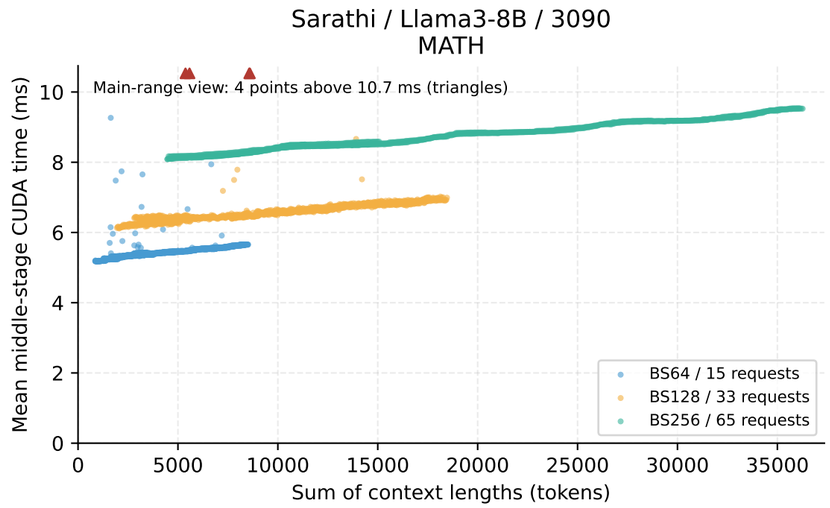
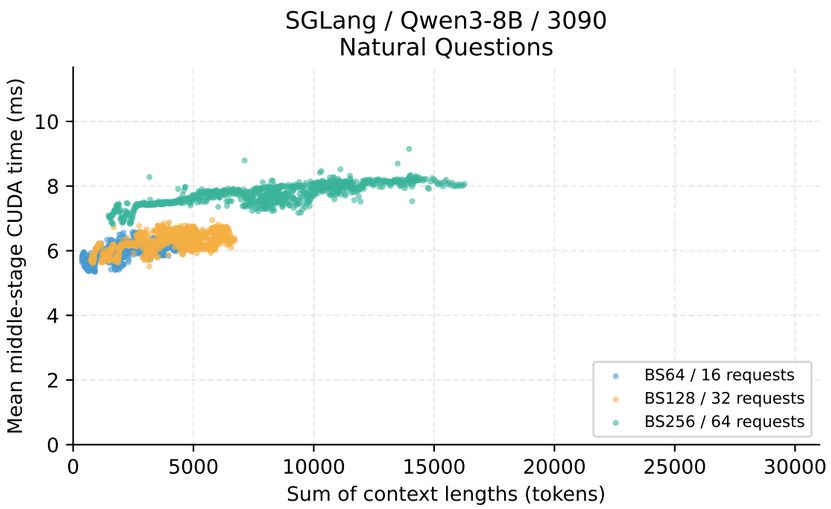
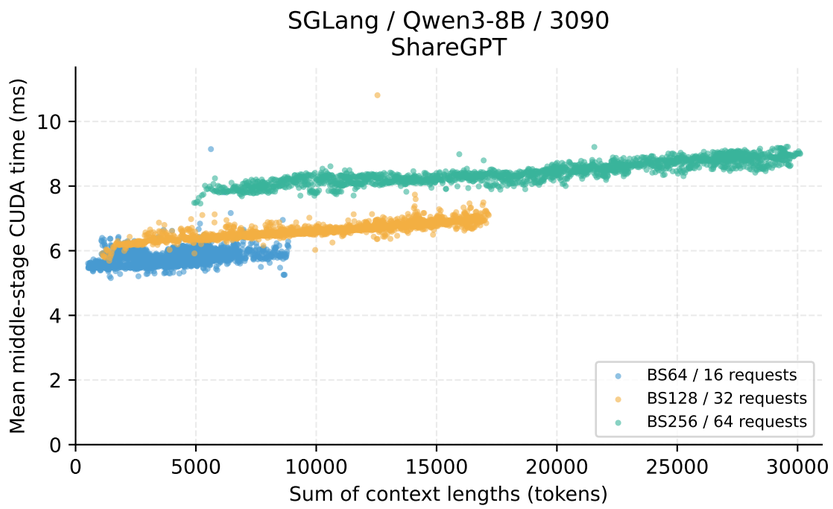
Auto, repeat 1 only | Fixed actual request count within each series

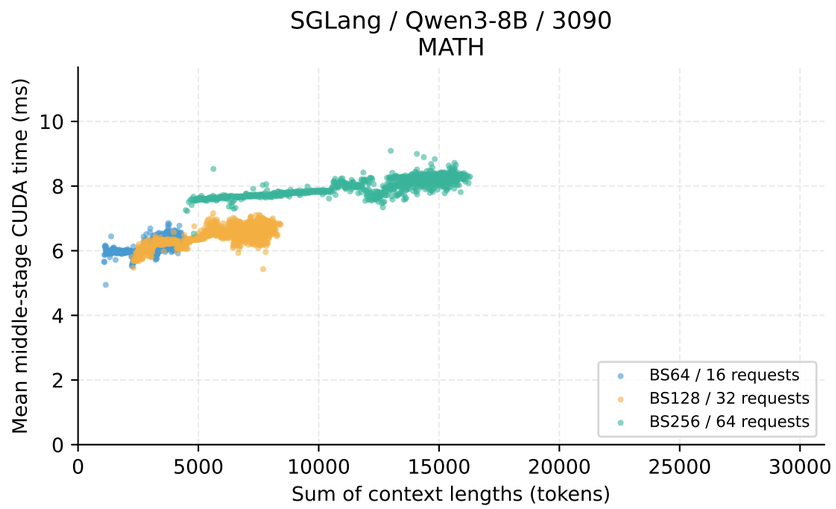


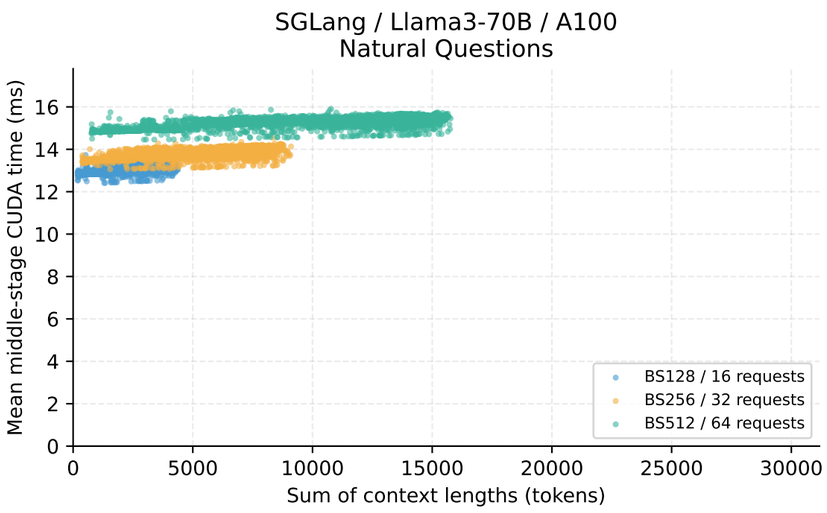


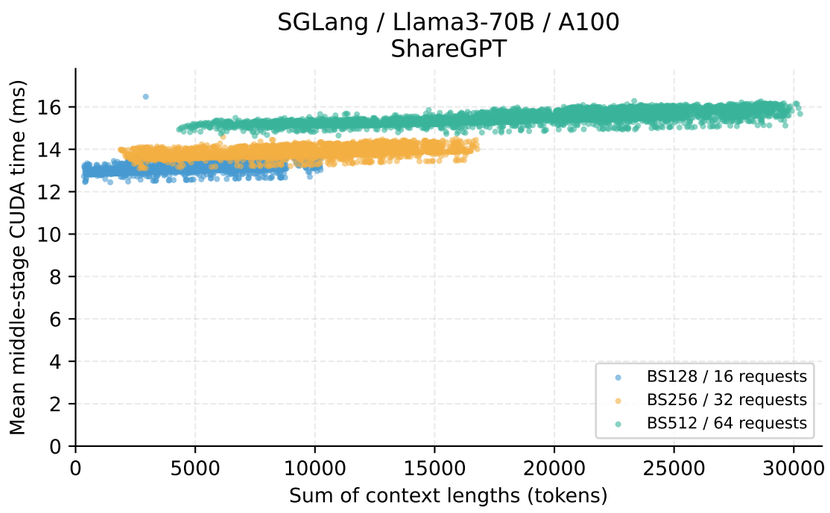


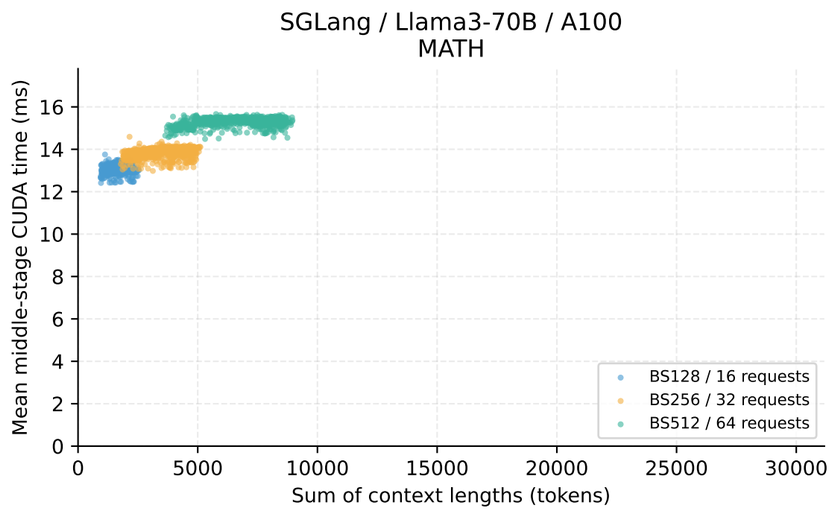












Equal x values averaged within each series; no time-based filtering or smoothing. Shared axes within each model row; y starts at zero.
Sarathi: most frequent pure-decode batch size per suite (see legends), middle stages 1–2. SGLang: BS/PP requests, middle stages 1–2 or 1–6.
Different models, hardware and workloads: this is a shape comparison, not a matched performance benchmark.