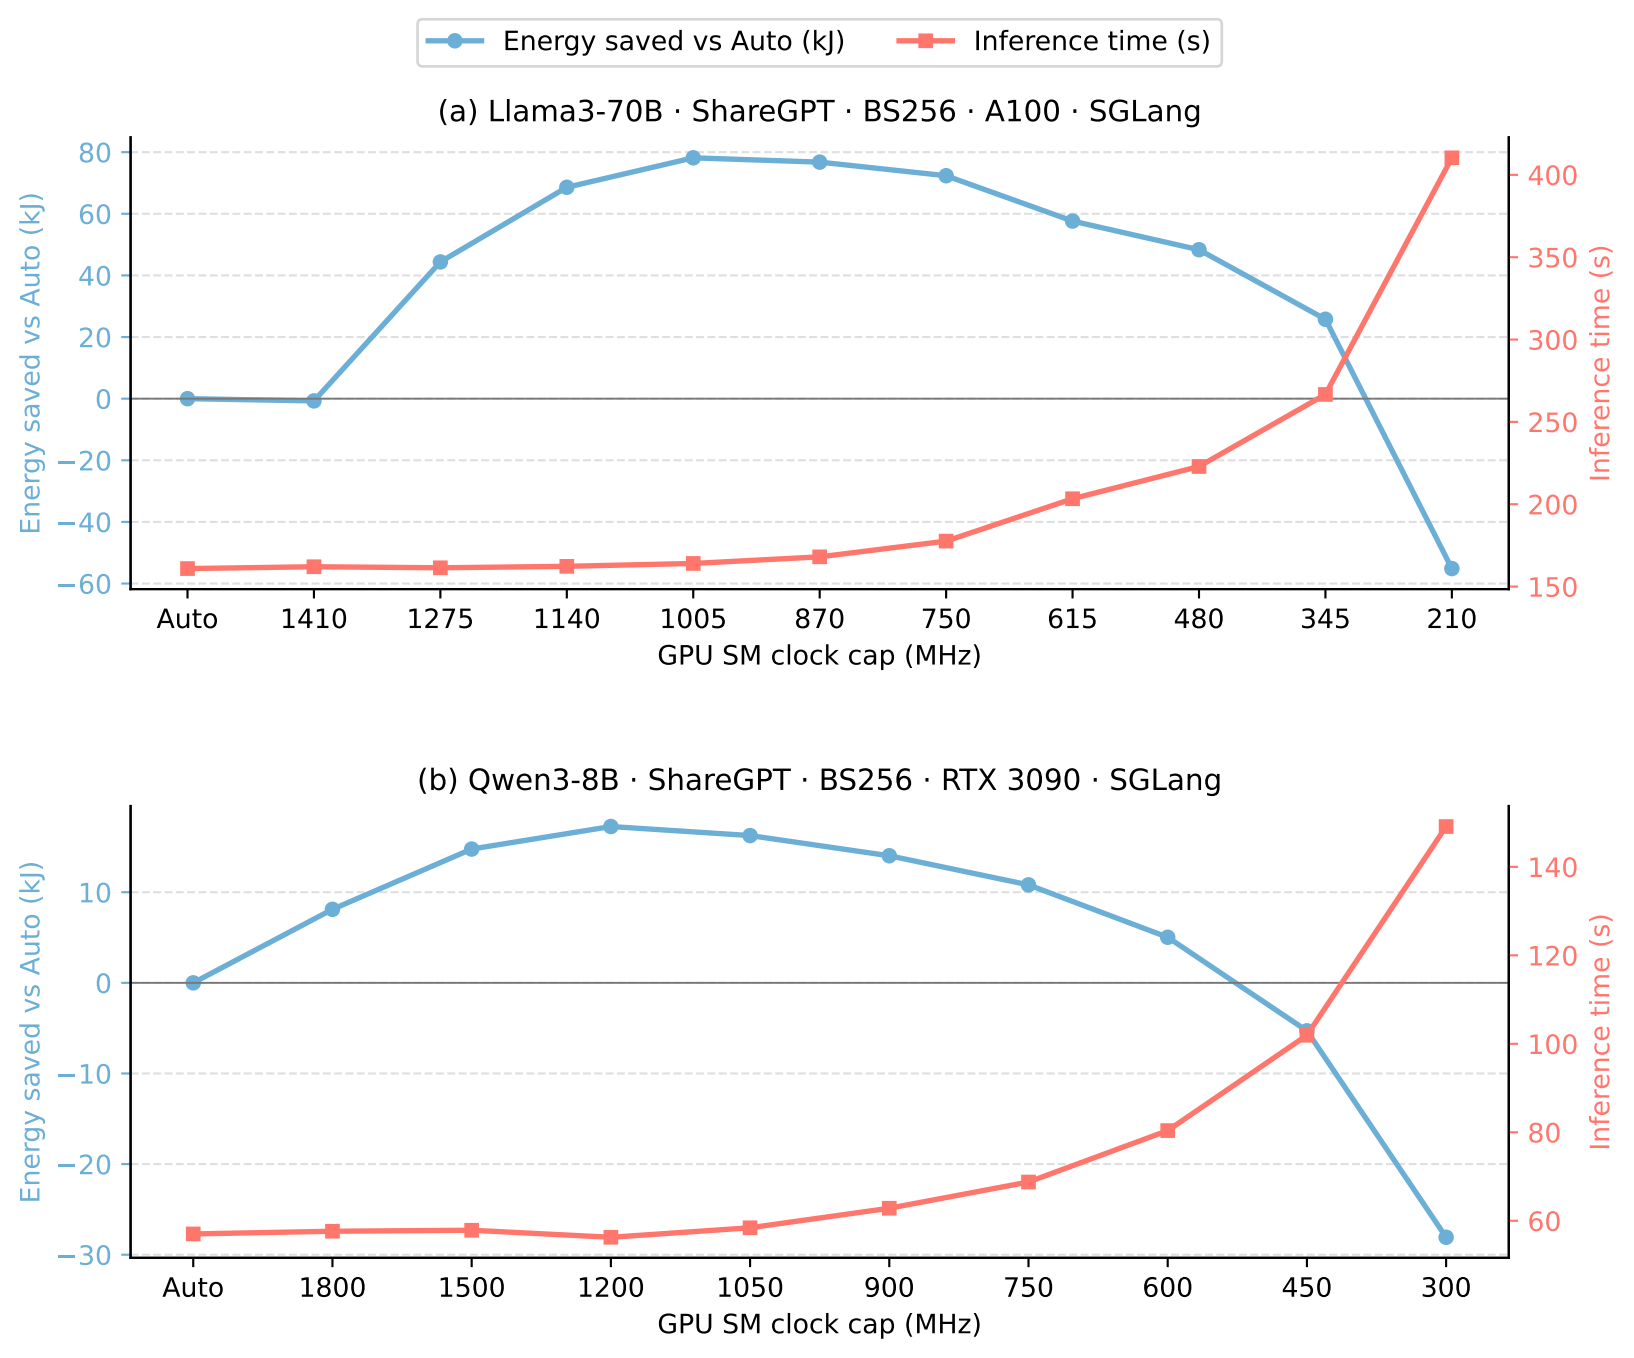
Energy saved vs Auto (kJ)
Inference time (s)
(a) Llama3-70B · ShareGPT · BS256 · A100 · SGLang
Energy saved vs Auto (kJ)
Inference time (s)
80
60
40
20
0
−20
−40
−60
400
350
300
250
200
150
Auto
1410
1275
1140
1005
870
750
615
480
345
210
GPU SM clock cap (MHz)
(b) Qwen3-8B · ShareGPT · BS256 · RTX 3090 · SGLang
Energy saved vs Auto (kJ)
Inference time (s)
10
0
−10
−20
−30
140
120
100
80
60
Auto
1800
1500
1200
1050
900
750
600
450
300
GPU SM clock cap (MHz)