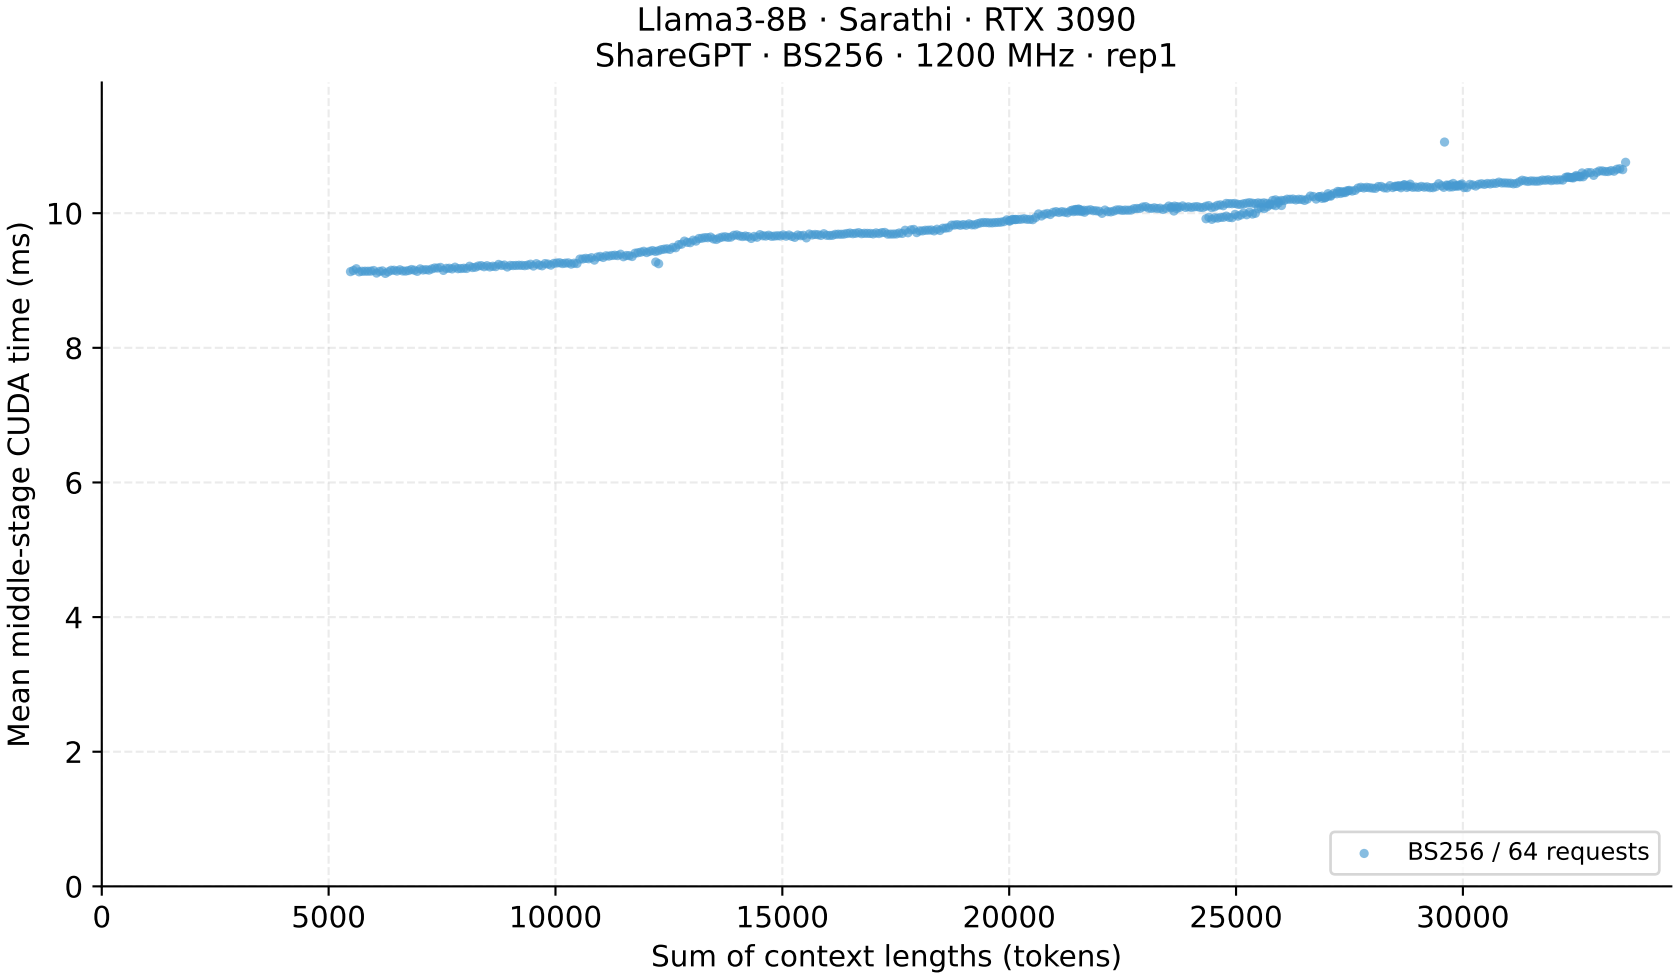
Llama3-8B · Sarathi · RTX 3090
ShareGPT · BS256 · 1200 MHz · rep1
Mean middle-stage CUDA time (ms)
0
2
4
6
8
10
0
5000
10000
15000
20000
25000
30000
Sum of context lengths (tokens)
BS256 / 64 requests