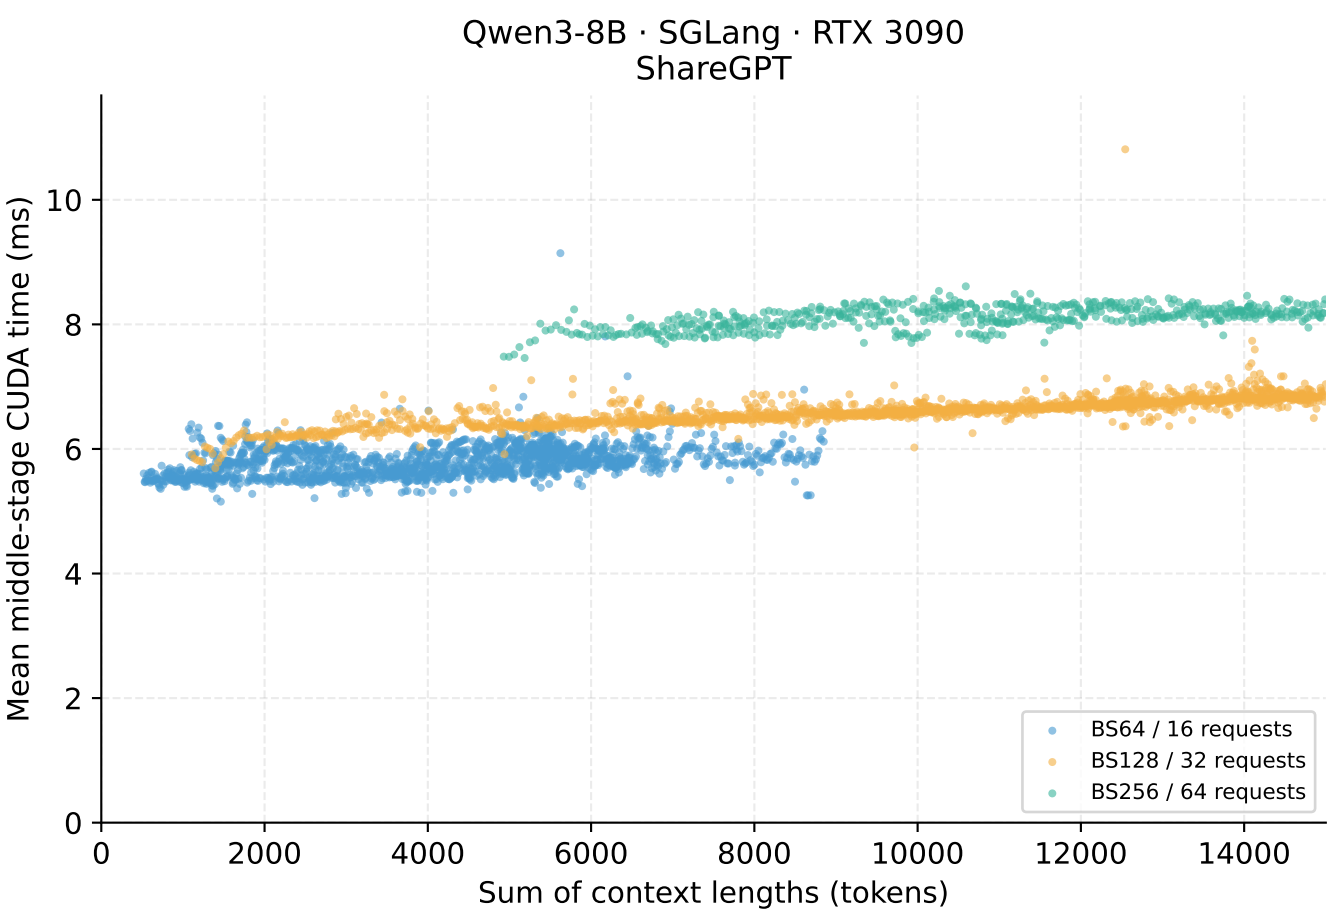
Qwen3-8B · SGLang · RTX 3090
ShareGPT
Mean middle-stage CUDA time (ms)
Sum of context lengths (tokens)
0
2
4
6
8
10
0
2000
4000
6000
8000
10000
12000
14000
BS64 / 16 requests
BS128 / 32 requests
BS256 / 64 requests