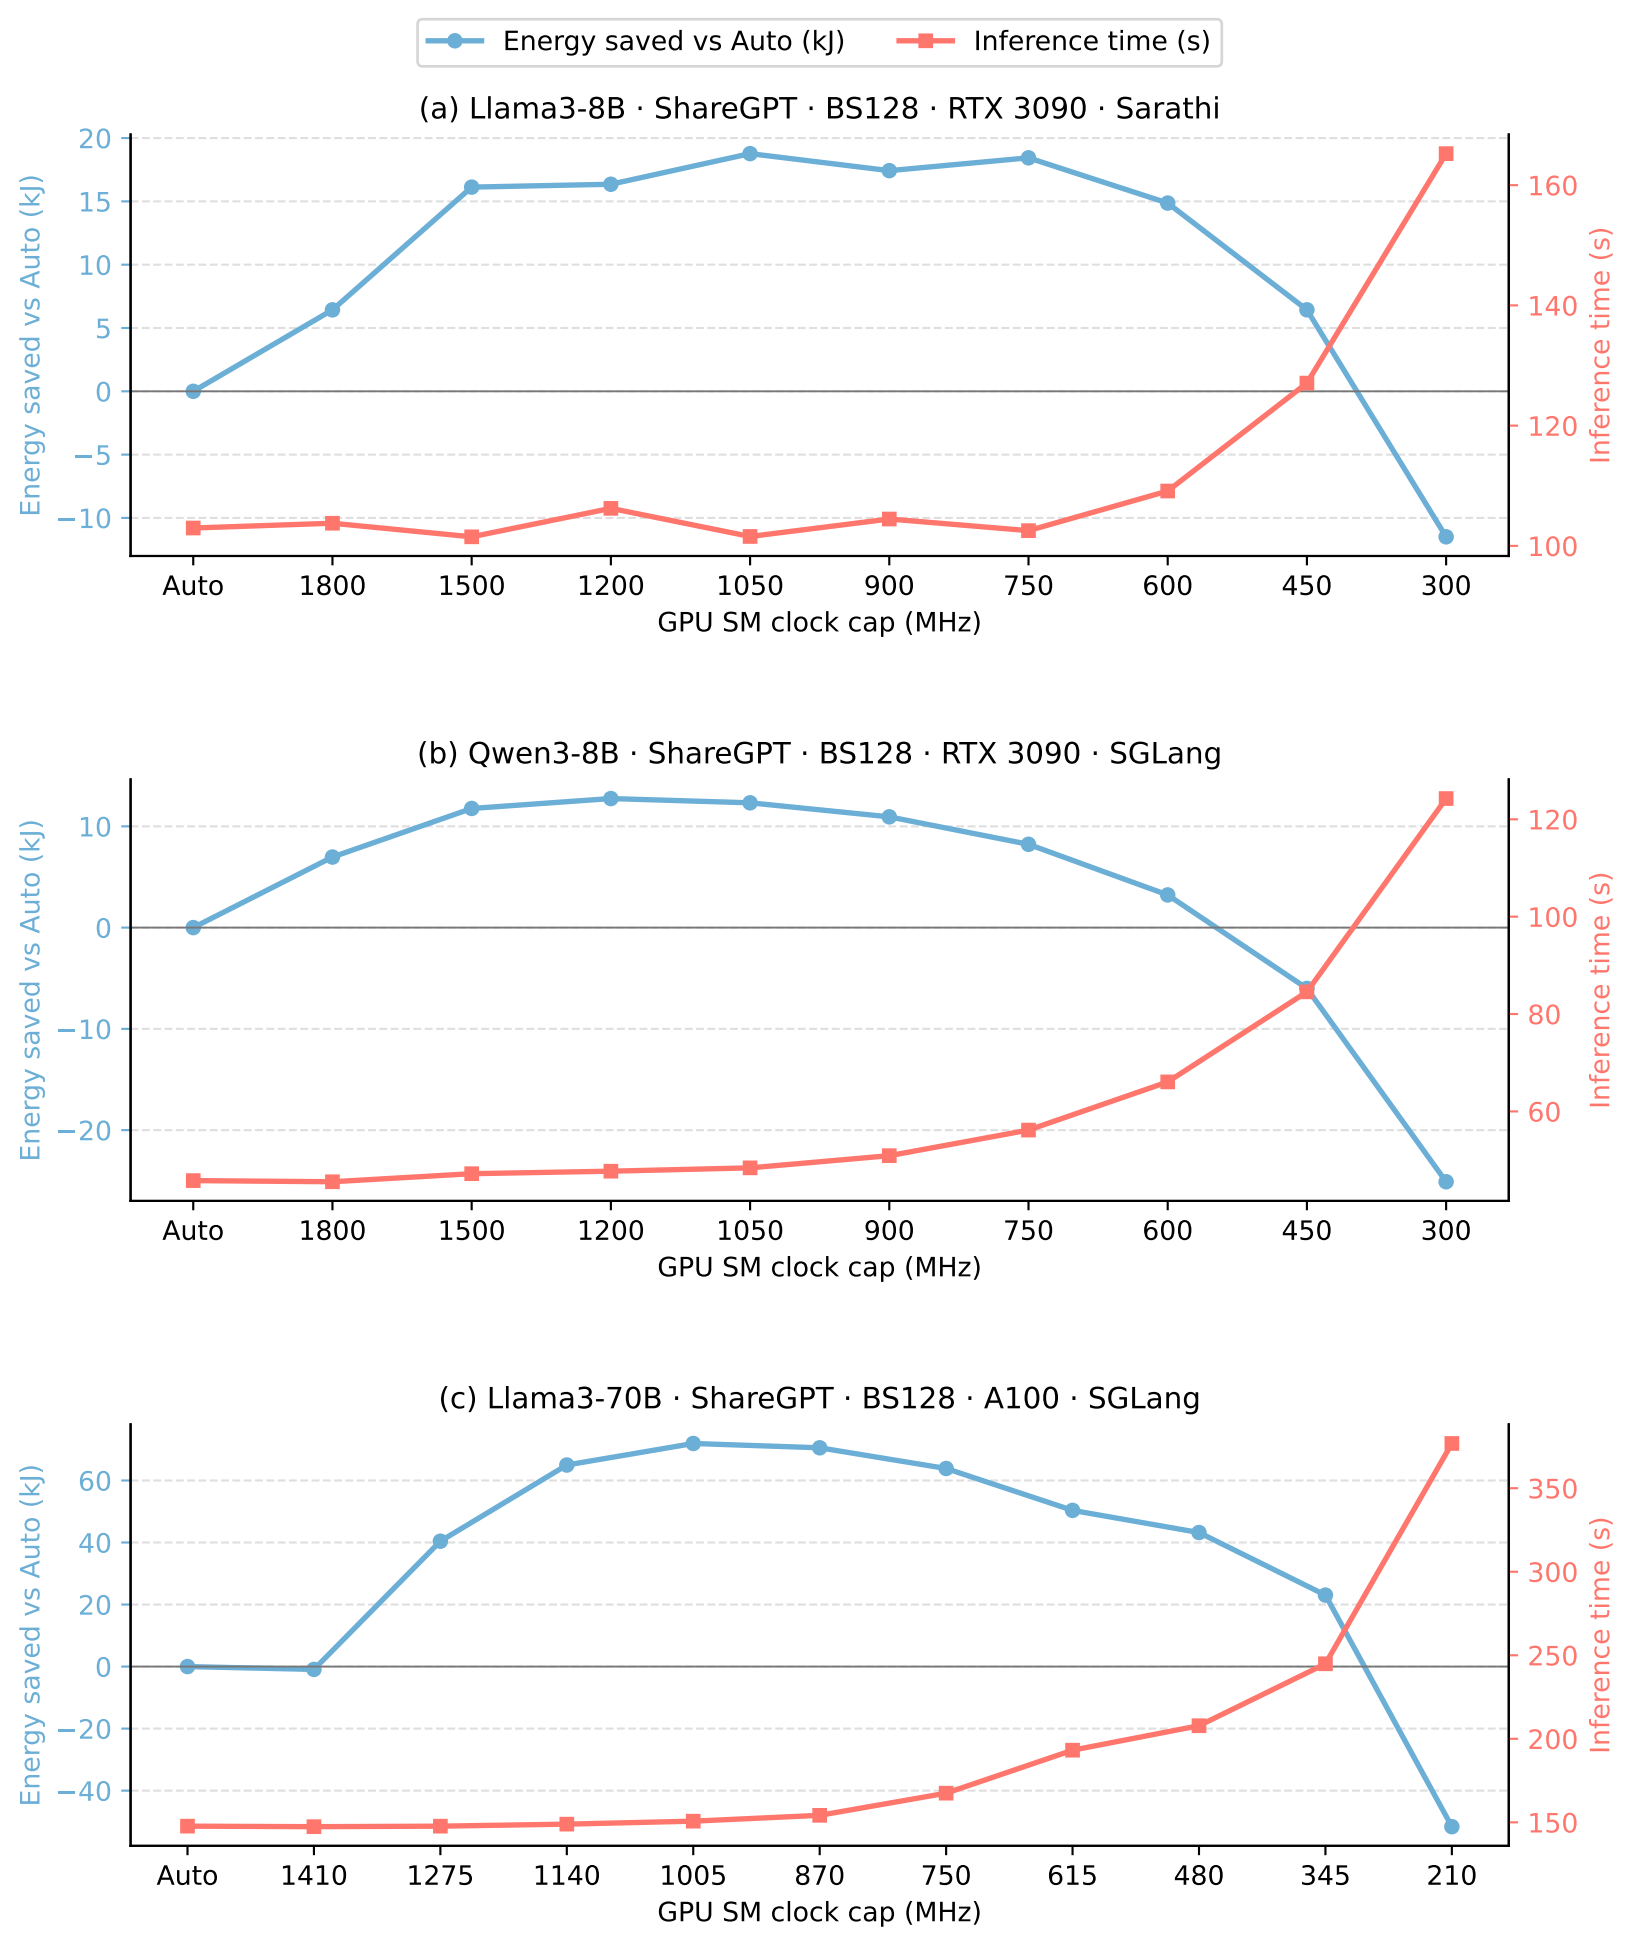
Energy saved vs Auto (kJ)
Inference time (s)
(a) Llama3-8B · ShareGPT · BS128 · RTX 3090 · Sarathi
Energy saved vs Auto (kJ)
Inference time (s)
GPU SM clock cap (MHz)
Auto 1800 1500 1200 1050 900 750 600 450 300
(b) Qwen3-8B · ShareGPT · BS128 · RTX 3090 · SGLang
Energy saved vs Auto (kJ)
Inference time (s)
GPU SM clock cap (MHz)
Auto 1800 1500 1200 1050 900 750 600 450 300
(c) Llama3-70B · ShareGPT · BS128 · A100 · SGLang
Energy saved vs Auto (kJ)
Inference time (s)
GPU SM clock cap (MHz)
Auto 1410 1275 1140 1005 870 750 615 480 345 210