Energy efficiency gain | Inference time change | Throughput decrease

ShareGPT
Llama3-8B · Sarathi

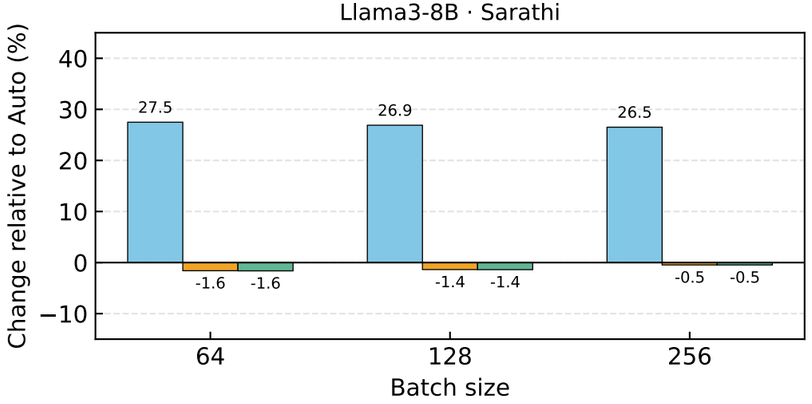

MATH
Llama3-8B · Sarathi

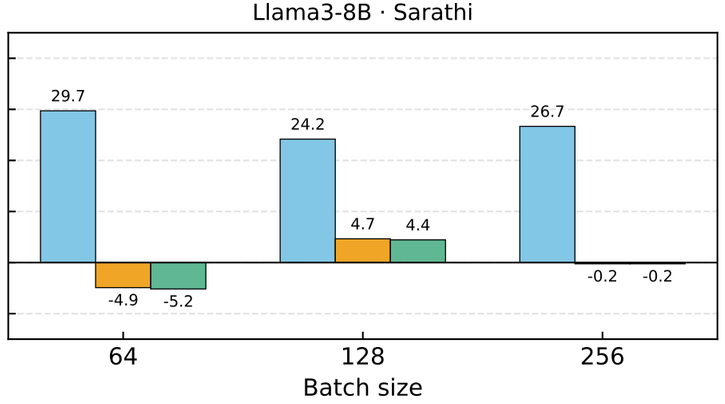

Natural Questions
Llama3-8B · Sarathi

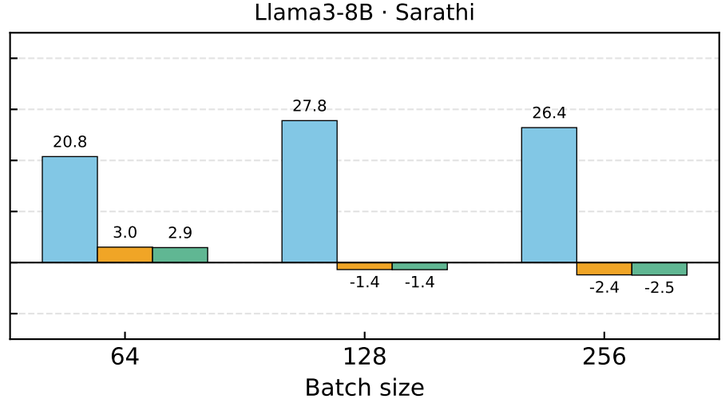

ShareGPT
Qwen3-8B · SGLang

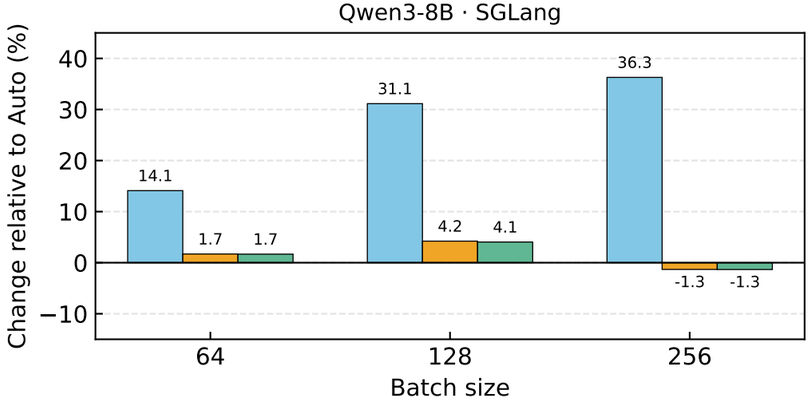

MATH
Qwen3-8B · SGLang

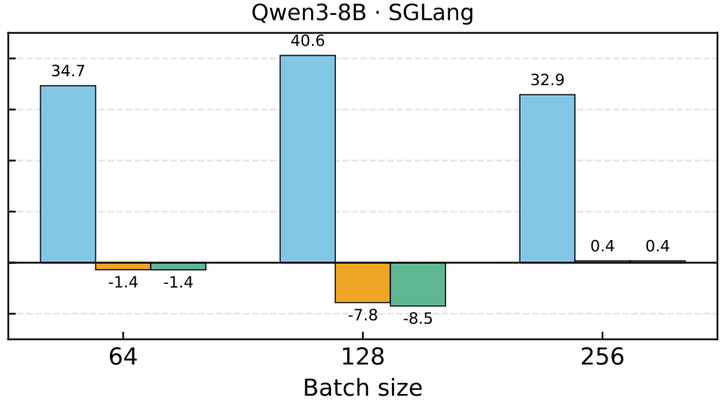

Natural Questions
Qwen3-8B · SGLang

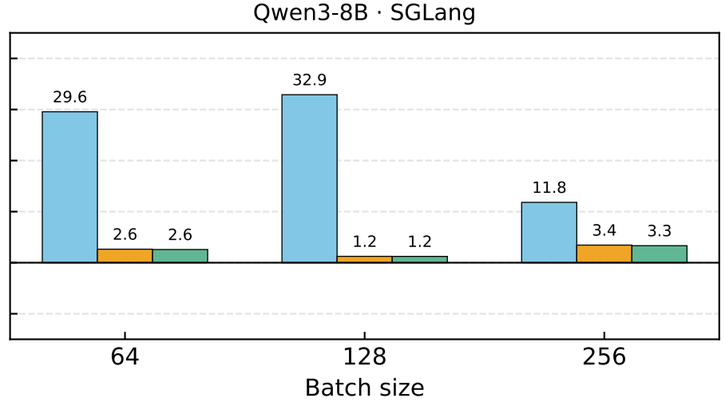

ShareGPT
Llama3-70B · SGLang · A100

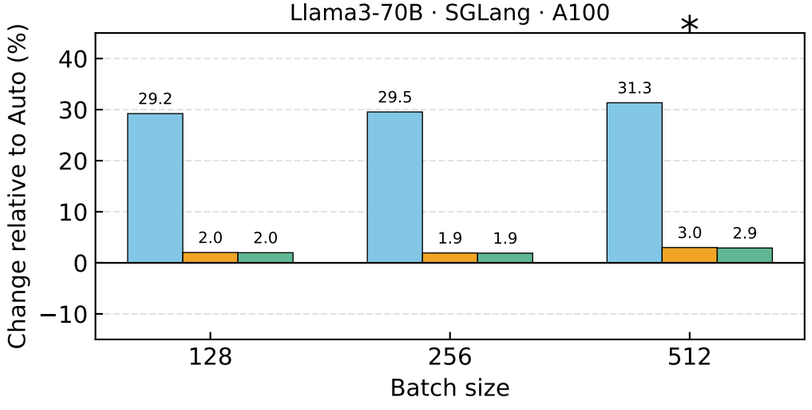

MATH
Llama3-70B · SGLang · A100

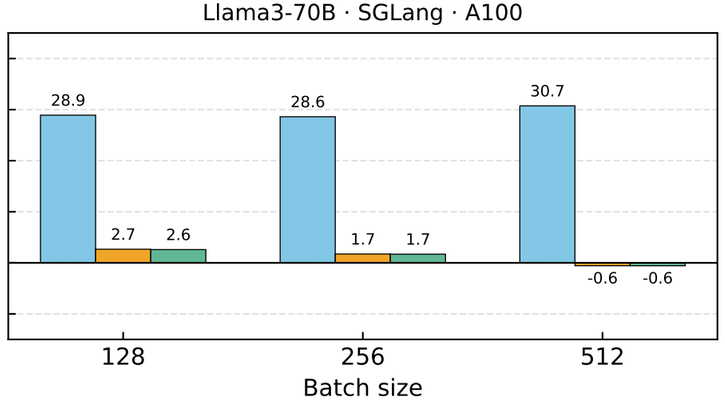

Natural Questions
Llama3-70B · SGLang · A100

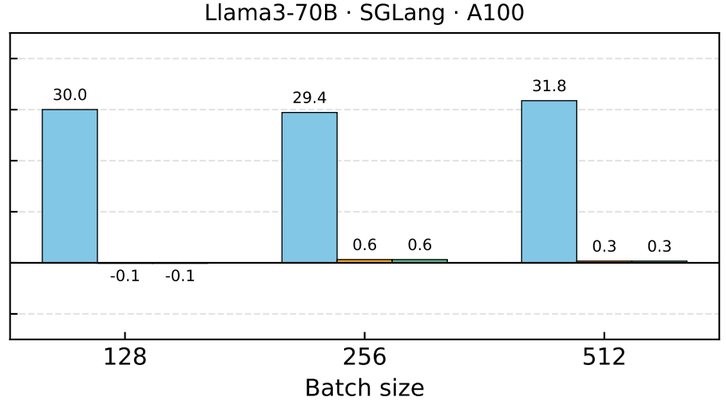

Lowest energy with ≤5% inference-time increase. A100 uses output2 (Stage 7 Auto). * ShareGPT BS512 aggregate is provisional.